



DeepSeek

DeepSeek is raising \$7B.
The efficiency lab just
went institutional.

swipe →



01

DeepSeek in 30 seconds

Chinese AI lab behind R1, an open model that matched top closed models at a fraction of training cost. The release rattled markets and changed how people thought about what open-source AI could do.



02

Their first funding round

About \$7 billion. Tencent puts in \$1.5B. CATL commits \$740M. Founder Liang Wenfeng contributes \$3B of his own capital. Valuation: \$52-59B post-money. Term sheets being signed as of June 4.



03

They preached efficiency

R1's whole story was doing more with less compute. Now they are on the same capital escalator as OpenAI and Anthropic. The lab that disrupted via frugality is raising like the labs it disrupted.



04

Scale still costs money

Even the efficiency lab needs billions to train next-gen models. The ceiling is higher than R1 showed. That is either humbling or clarifying, depending on where you sit.



05

Watch the license terms

DeepSeek has stayed open so far. Institutional investors typically want returns. Whether the open weights continue after this round matters a lot for anyone building on them.



Open source or institutional control?

The next models DeepSeek ships may be the answer.

[Follow Ahmed Raza for builder takes on AI](#)