



NVIDIA

Nvidia's 550B open model: what it means for builders

swipe →



01

550B total. 55B active.

MoE with 10:1 sparsity. Only 55B parameters fire per token. That's why you get 300+ tokens per second without giving up depth.



02

1 million token context

Not a benchmark footnote. Agents doing multi-step planning over large codebases or documents can hold more in memory than any US open model before this.



03

Open weights, real license

NVIDIA Open Model License. Commercial use permitted. Weights on HuggingFace. API on NVIDIA NIM and OpenRouter. Fully self-hostable.



04

The honest gap

Kimi K2.6 scores 54 on Artificial Analysis' Intelligence Index. Nemotron Ultra scores 48. Best US open model. Not the global leader.



05

The infrastructure reality

300GB+ VRAM to self-host. A100 or H100 clusters minimum. Most builders are better off on the NIM or OpenRouter API.



What changes for you?

A US-made open frontier model with commercial rights and 1M context. Does model provenance actually change how you architect?

Follow Ahmed Raza for builder takes on AI